

```
In [196... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [197... housing=pd.read_csv(r"D:\PC\Data Science\data sets\Housing\housing.csv")
```

```
In [198... housing.head()
```

```
Out[198... longitude latitude housing_median_age total_rooms total_bedrooms population households
```

0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0

```
In [199... housing.columns
```

```
Out[199... Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
        'total_bedrooms', 'population', 'households', 'median_income',
        'median_house_value', 'ocean_proximity'],
        dtype='object')
```

```
In [200... housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null float64
1   latitude               20640 non-null float64
2   housing_median_age     20640 non-null float64
3   total_rooms            20640 non-null float64
4   total_bedrooms         20433 non-null float64
5   population             20640 non-null float64
6   households             20640 non-null float64
7   median_income          20640 non-null float64
8   median_house_value     20640 non-null float64
9   ocean_proximity        20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
In [201... housing["ocean_proximity"].value_counts()
```

```
Out[201... <1H OCEAN      9136
INLAND         6551
NEAR OCEAN     2658
NEAR BAY       2290
ISLAND          5
Name: ocean_proximity, dtype: int64
```

```
In [202... housing.describe()
```

```
Out[202... longitude latitude housing_median_age total_rooms total_bedrooms population
```

count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
<b>std</b>	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122
<b>min</b>	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000
<b>25%</b>	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000
<b>50%</b>	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000
<b>75%</b>	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000
<b>max</b>	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000



In [203...] `from sklearn.model_selection import train_test_split`


In [204...] `train_set, test_set= train_test_split(housing, test_size=0.2, random_state=42)`

In [205...] `housing["income_cat"] = pd.cut(housing["median_income"],  
bins=[0., 1.5, 3.0, 4.5, 6., np.inf],  
labels=[1, 2, 3, 4, 5])`

In [206...] `housing_labels = train_set["median_house_value"].copy()`


In [207...] `housing["income_cat"].hist()`

Out[207...] `<AxesSubplot:>`




In [208...] `housing.plot(kind="scatter", x="longitude", y="latitude")`

Out[208...] `<AxesSubplot:xlabel='longitude', ylabel='latitude'>`



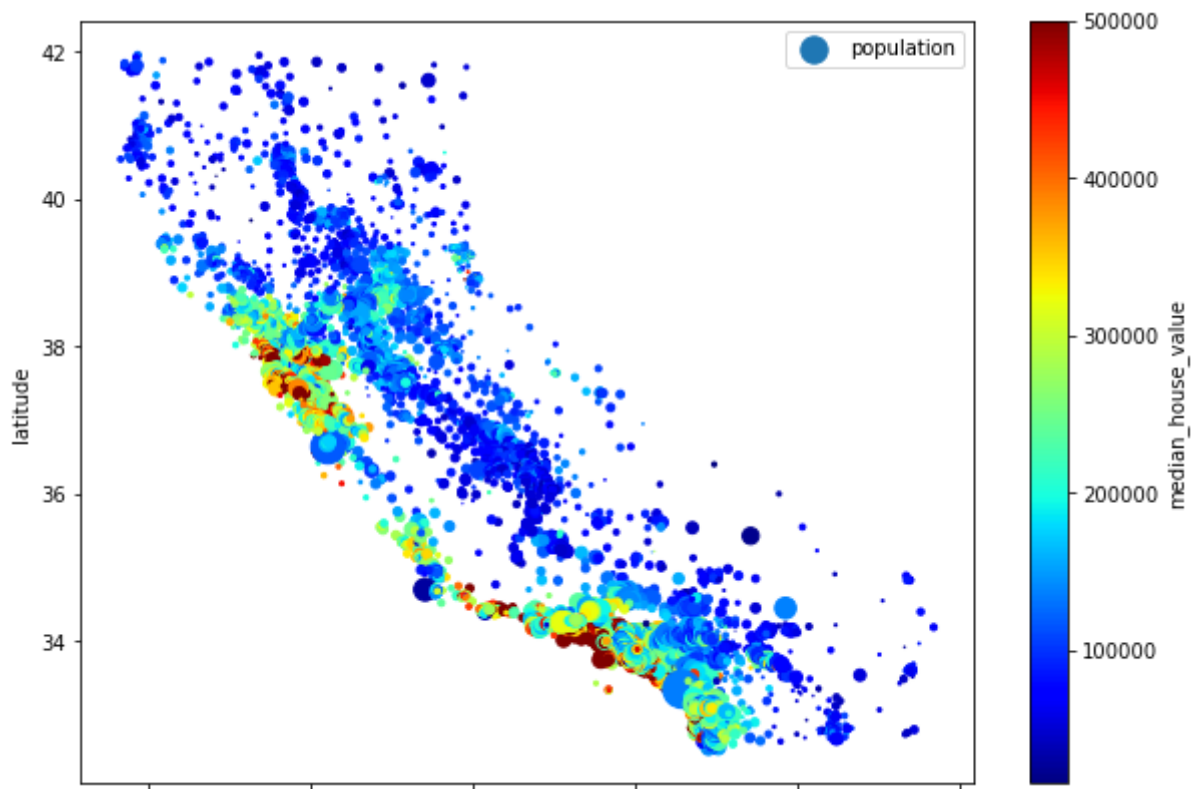
```
In [209...] housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)
```

```
Out[209...] <AxesSubplot:xlabel='longitude', ylabel='latitude'>
```



```
In [210...] housing.plot(kind="scatter", x="longitude", y="latitude",  
                          s=housing["population"]/100, label="population", figsize=(10,7),  
                          c="median_house_value", cmap=plt.get_cmap("jet"), colorbar=True,)  
plt.legend()
```

```
Out[210...] <matplotlib.legend.Legend at 0x2059fe222e0>
```



```
In [211...] corr_matrix= housing.corr()
```

```
In [212...] corr_matrix["median_house_value"].sort_values(ascending=False)
```

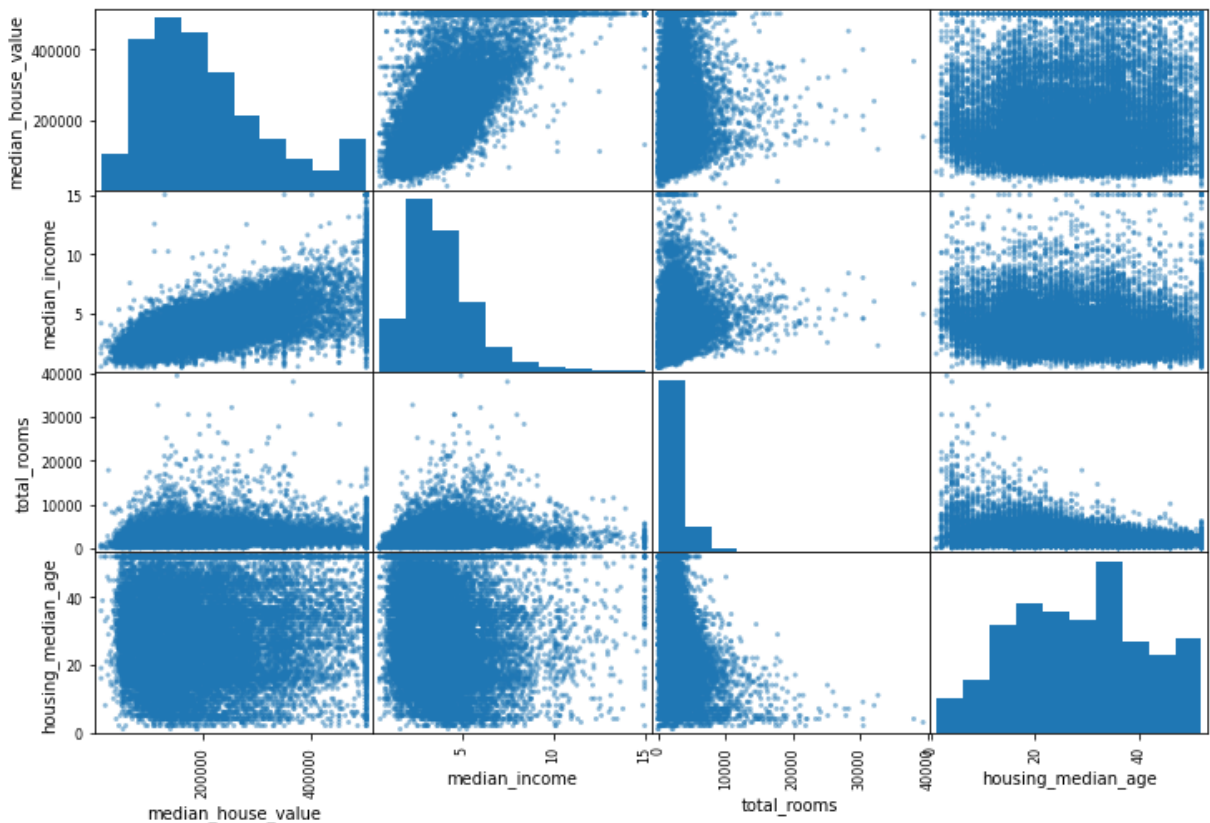
```
Out[212...] median_house_value    1.000000
median_income         0.688075
total_rooms           0.134153
housing_median_age    0.105623
households            0.065843
total_bedrooms        0.049686
population            -0.024650
longitude             -0.045967
latitude              -0.144160
Name: median_house_value, dtype: float64
```

```
In [213...] from pandas.plotting import scatter_matrix
```

```
In [214...] attributes = ["median_house_value", "median_income", "total_rooms", "housing_median_
```

```
In [215...] scatter_matrix(housing[attributes], figsize=(12,8))
```

```
Out[215...] array([[<AxesSubplot:xlabel='median_house_value', ylabel='median_house_value'>,
<AxesSubplot:xlabel='median_income', ylabel='median_house_value'>,
<AxesSubplot:xlabel='total_rooms', ylabel='median_house_value'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='median_house_value'>],
[<AxesSubplot:xlabel='median_house_value', ylabel='median_income'>,
<AxesSubplot:xlabel='median_income', ylabel='median_income'>,
<AxesSubplot:xlabel='total_rooms', ylabel='median_income'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='median_income'>],
[<AxesSubplot:xlabel='median_house_value', ylabel='total_rooms'>,
<AxesSubplot:xlabel='median_income', ylabel='total_rooms'>,
<AxesSubplot:xlabel='total_rooms', ylabel='total_rooms'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='total_rooms'>],
[<AxesSubplot:xlabel='median_house_value', ylabel='housing_median_age'>,
<AxesSubplot:xlabel='median_income', ylabel='housing_median_age'>,
<AxesSubplot:xlabel='total_rooms', ylabel='housing_median_age'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='housing_median_age'>]],
dtype=object)
```



```
In [216...] median= train_set["total_bedrooms"].median()
```

```
In [217...] train_set["total_bedrooms"].fillna(median, inplace=True)
```

<ipython-input-217-9ba39919f72d>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
 train\_set["total\_bedrooms"].fillna(median, inplace=True)

```
In [248...] train_set.columns
```

```
Out[248...] Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',  
              'total_bedrooms', 'population', 'households', 'median_income',  
              'median_house_value', 'ocean_proximity'],  
              dtype='object')
```

```
In [249...] X_train_set=train_set.drop("median_house_value", axis=1)
```

```
In [250...] from sklearn.impute import SimpleImputer
```

```
In [251...] housing_num=X_train_set.drop("ocean_proximity", axis=1)
```

```
In [252...] imputer=SimpleImputer(strategy="median")
```

```
In [253...] imputer.fit(housing_num)
```

```
Out[253...] SimpleImputer(strategy='median')
```

```
In [254...] imputer.statistics_
```

```
Out[254...] array([-118.51 ,  34.26 ,  29.    , 2129.    ,  437.    , 1167.    ,  
                410.    ,  3.5458])
```

```
In [255...] housing_num.median().values
```

```
Out[255...] array([-118.51 ,  34.26 ,  29.    , 2129.    ,  437.    , 1167.    ,
          410.    ,  3.5458])
```

```
In [256...] X=imputer.transform(housing_num)
```

```
In [257...] housing_tr=pd.DataFrame(X, columns=housing_num.columns)
```

```
In [258...] housing_cat=train_set[["ocean_proximity"]]
```

```
In [259...] housing_cat.head()
```

```
Out[259...]
ocean_proximity
14196    NEAR OCEAN
8267     NEAR OCEAN
17445    NEAR OCEAN
14265    NEAR OCEAN
2271         INLAND
```

```
In [260...] from sklearn.preprocessing import OrdinalEncoder
```

```
In [261...] ordinal_encoder=OrdinalEncoder()
```

```
In [262...] housing_cat_encoded= ordinal_encoder.fit_transform(housing_cat)
```

```
In [263...] housing_cat_encoded[:10]
```

```
Out[263...] array([[4.],
          [4.],
          [4.],
          [4.],
          [1.],
          [0.],
          [0.],
          [3.],
          [0.],
          [0.]])
```

```
In [264...] ordinal_encoder.categories_
```

```
Out[264...] [array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],
      dtype=object)]
```

```
In [265...] from sklearn.preprocessing import OneHotEncoder
```

```
In [266...] cat_encoder=OneHotEncoder()
```

```
In [267...] housing_cat_1hot= cat_encoder.fit_transform(housing_cat)
```

```
In [268...] housing_cat_1hot
```

```
Out[268...] <16512x5 sparse matrix of type '<class 'numpy.float64'>'
      with 16512 stored elements in Compressed Sparse Row format>
```

```
In [269...] housing_cat_1hot.toarray()
```

```
Out[269...] array([[0., 0., 0., 0., 1.],
        [0., 0., 0., 0., 1.],
        [0., 0., 0., 0., 1.],
        ...,
        [1., 0., 0., 0., 0.],
        [1., 0., 0., 0., 0.],
        [0., 0., 0., 1., 0.]])
```

```
In [270...] cat_encoder.categories_
```

```
Out[270...] [array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],
        dtype=object)]
```

```
In [271...] from sklearn.base import BaseEstimator, TransformerMixin
```

```
In [272...] from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
```

```
In [273...] num_pipeline= Pipeline([
        ('imputer', SimpleImputer(strategy="median")),
        ('std_scaler', StandardScaler()),
    ])
```

```
In [274...] housing_num_tr=num_pipeline.fit_transform(housing_num)
```

```
In [275...] from sklearn.compose import ColumnTransformer
```

```
In [276...] num_attribs= list(housing_num)
cat_attribs= ["ocean_proximity"]
```

```
In [345...] full_pipeline= ColumnTransformer([
        ("num", num_pipeline, num_attribs),
        ("cat", OneHotEncoder(), cat_attribs),
    ])
```

```
In [351...] train_set.head()
```

```
Out[351...]      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  population  househo
14196    -117.03    32.71             33.0         3126.0           627.0       2300.0         62
8267     -118.16    33.77             49.0         3382.0           787.0       1314.0         75
17445    -120.48    34.66              4.0         1897.0           331.0        915.0         33
14265    -117.11    32.69             36.0         1421.0           367.0       1418.0         35
2271     -119.80    36.78             43.0         2382.0           431.0        874.0         38
```



```
In [347...] housing_prepared= full_pipeline.fit_transform(train_set)
```

```
In [279...] from sklearn.linear_model import LinearRegression
```

```
In [280...] lin_reg= LinearRegression()
```

```
In [281...] lin_reg.fit(housing_prepared, housing_labels)
```

```
Out[281...] LinearRegression()
```

```
In [282... from sklearn.metrics import mean_squared_error
```

```
In [283... housing_predictions= lin_reg.predict(housing_prepared)
```

```
In [284... lin_mse= mean_squared_error(housing_labels, housing_predictions)
```

```
In [285... lin_rmse= np.sqrt(lin_mse)
```

```
In [286... lin_rmse
```

```
Out[286... 68433.93736666226
```

```
In [287... from sklearn.metrics import r2_score
```

```
In [288... r2_score(housing_labels, housing_predictions)
```

```
Out[288... 0.6496648627123223
```

```
In [289... from sklearn.tree import DecisionTreeRegressor
```

```
In [290... tree_reg= DecisionTreeRegressor()
```

```
In [291... tree_reg.fit(housing_prepared, housing_labels)
```

```
Out[291... DecisionTreeRegressor()
```

```
In [292... housing_predictions= tree_reg.predict(housing_prepared)
```

```
In [293... tree_mse=mean_squared_error(housing_labels, housing_predictions)
```

```
In [294... tree_rmse=np.sqrt(tree_mse)
```

```
In [295... tree_rmse
```

```
Out[295... 0.0
```

```
In [296... from sklearn.model_selection import cross_val_score
```

```
In [301... scores= cross_val_score(tree_reg, housing_prepared, housing_labels,  
                           scoring="neg_mean_squared_error", cv=10)
```

```
In [302... tree_rmse_scores= np.sqrt(-scores)
```

```
In [303... def display_scores(scores):  
    print("Scores", scores)  
    print("Mean", scores.mean())  
    print("SD", scores.std())
```

```
In [304... display_scores(tree_rmse_scores)
```

```
Scores [66276.3541399  67726.07353866 67990.65131197 72371.60417774  
        66988.98321269 66475.391629   62779.49161703 69994.77542875  
        70474.54007364 67118.74896264]  
Mean 67819.66140920234  
SD 2516.170812019729
```

```
In [311... from sklearn.ensemble import RandomForestRegressor
```



```
In [312... forest_reg= RandomForestRegressor()
```

```
In [313... forest_reg.fit(housing_prepared, housing_labels)
```

```
Out[313... RandomForestRegressor()
```

```
In [314... housing_predictions= forest_reg.predict(housing_prepared)
```

```
In [315... forest_mse=mean_squared_error(housing_labels, housing_predictions)
```

```
In [316... forest_rmse=np.sqrt(forest_mse)
```

```
In [317... forest_rmse
```

```
Out[317... 18094.098051516652
```

```
In [318... scores= cross_val_score(forest_reg, housing_prepared, housing_labels,  
                           scoring="neg_mean_squared_error", cv=10)
```

```
In [319... forest_rmse_scores= np.sqrt(-scores)
```

```
In [320... display_scores(forest_rmse_scores)
```

```
Scores [46925.29728707 50633.14879565 47809.58201204 50158.44551746  
        49777.14297837 46489.55801507 45669.27129373 50812.68309854  
        49381.10350237 49709.55915788]  
Mean 48736.579165819196  
SD 1759.866755098333
```

```
In [321... r2_score(housing_labels, housing_predictions)
```

```
Out[321... 0.9755085487322805
```

```
In [322... from sklearn.model_selection import GridSearchCV
```

```
In [325... param_grid=[  
    {'n_estimators':[3,10,30], 'max_features':[2,4,6,8]},  
    {'bootstrap':[False], 'n_estimators':[3,10], 'max_features':[2,3,4]},  
    ]
```

```
In [326... forest_reg= RandomForestRegressor()
```

```
In [327... grid_search=GridSearchCV(forest_reg, param_grid, cv=5,  
                           scoring='neg_mean_squared_error',  
                           return_train_score=True)
```

```
In [328... grid_search.fit(housing_prepared, housing_labels)
```

```
Out[328... GridSearchCV(cv=5, estimator=RandomForestRegressor(),  
                param_grid=[{'max_features': [2, 4, 6, 8],  
                             'n_estimators': [3, 10, 30]},  
                {'bootstrap': [False], 'max_features': [2, 3, 4],  
                             'n_estimators': [3, 10]}],  
                return_train_score=True, scoring='neg_mean_squared_error')
```

```
In [329... grid_search.best_params_
```

```
Out[329... {'max_features': 8, 'n_estimators': 30}
```

In [330...] `grid_search.best_estimator_`

Out[330...] `RandomForestRegressor(max_features=8, n_estimators=30)`

In [331...] `cvres= grid_search.cv_results_`

In [333...] `for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):  
 print(np.sqrt(-mean_score), params)`

```
62692.4654345404 {'max_features': 2, 'n_estimators': 3}
55069.50733235221 {'max_features': 2, 'n_estimators': 10}
52110.59077103753 {'max_features': 2, 'n_estimators': 30}
58955.342263822124 {'max_features': 4, 'n_estimators': 3}
52297.74217779409 {'max_features': 4, 'n_estimators': 10}
50181.75344023495 {'max_features': 4, 'n_estimators': 30}
58999.39400349141 {'max_features': 6, 'n_estimators': 3}
52285.04447774024 {'max_features': 6, 'n_estimators': 10}
49858.03618233968 {'max_features': 6, 'n_estimators': 30}
58826.82351294995 {'max_features': 8, 'n_estimators': 3}
51469.744462704686 {'max_features': 8, 'n_estimators': 10}
49541.36348219184 {'max_features': 8, 'n_estimators': 30}
62069.12781237627 {'bootstrap': False, 'max_features': 2, 'n_estimators': 3}
54450.758248770755 {'bootstrap': False, 'max_features': 2, 'n_estimators': 10}
59410.50337615474 {'bootstrap': False, 'max_features': 3, 'n_estimators': 3}
52428.87435495705 {'bootstrap': False, 'max_features': 3, 'n_estimators': 10}
58138.52190844188 {'bootstrap': False, 'max_features': 4, 'n_estimators': 3}
51251.93030127868 {'bootstrap': False, 'max_features': 4, 'n_estimators': 10}
```

In [334...] `final_model= grid_search.best_estimator_`

In [360...] `X_test=test_set.drop("median_house_value", axis=1)`

In [342...] `y_test= test_set["median_house_value"].copy()`

In [354...] `test_set.head()`

Out[354...] 

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	househo
<b>20046</b>	-119.01	36.06	25.0	1505.0	NaN	1392.0	35
<b>3024</b>	-119.46	35.14	30.0	2943.0	NaN	1565.0	58
<b>15663</b>	-122.44	37.80	52.0	3830.0	NaN	1310.0	96
<b>20484</b>	-118.72	34.28	17.0	3051.0	NaN	1705.0	49
<b>9814</b>	-121.93	36.62	34.0	2351.0	NaN	1063.0	42



In [361...] `test_set.columns`

Out[361...] `Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',  
 'total_bedrooms', 'population', 'households', 'median_income',  
 'median_house_value', 'ocean_proximity'],  
 dtype='object')`

In [363...] `X_test_prepared= full_pipeline.transform(test_set)`

In [364...] `final_predictions= final_model.predict(X_test_prepared)`

In [365...] `final_mse= mean_squared_error(y_test, final_predictions)`

```
In [366... final_rmse= np.sqrt(final_mse)
```

```
In [367... final_rmse
```

```
Out[367... 48586.34629402553
```

```
In [ ]:
```